

Exercise Session XI: DNS Security

Network Security

Matteo Scarlata

December 3rd, 2020

ETH Zurich

Context: Domain Name System

DNS: Origins

```
HOST : 192.5.89.50 : KENTARUS.ARPA,KENTARUS : SUN-2/120 : UNIX : TCP/TELNET,TCP/SMTP,TCP/FTP,ICMP,UDP :
HOST : 192.12.59.50 : MCCDB-ERATOSTHENES.ARPA,MCCDB-ERATOSTHENES,ERATOSTHENES : SUN-2/170 : UNIX :
TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.51 : TAUCETI.ARPA,TAUCETI : SUN-2/120 : UNIX : TCP/TELNET,TCP/SMTP,TCP/FTP,ICMP,UDP :
HOST : 192.12.59.51 : MCCDB-ACKERMANN.ARPA,MCCDB-ACKERMANN,ACKERMANN : SUN-2/50 : UNIX :
TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.12.59.52 : MCCDB-HILBERT.ARPA,MCCDB-HILBERT,HILBERT : SUN-2/50 : UNIX :
TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.12.59.53 : MCCDB-SUPPES.ARPA,MCCDB-SUPPES,SUPPES : SUN-2/50 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.12.59.54 : MCCDB-ZERMELO.ARPA,MCCDB-ZERMELO,ZERMELO : SUN-2/50 : UNIX :
TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.2.55, 192.12.12.1 : WISC-SPOOL.ARPA,WISC-SPOOL,LOOPS,NEEDLE : VAX-11/750 : UNIX :
TCP/TELNET,TCP/FTP,TCP/SMTP,TCP/ECHO,TCP/FINGER,ICMP :
HOST : 192.12.59.55 : MCCDB-TARTAGLIA.ARPA,MCCDB-TARTAGLIA,TARTAGLIA : SUN-2/50 : UNIX :
TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.12.59.56 : MCCDB-VIETE.ARPA,MCCDB-VIETE,VIETE : SUN-2/50 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.12.59.57 : MCCDB-LEGENDRE.ARPA,MCCDB-LEGENDRE,LEGENDRE : SUN-2/50 : UNIX :
TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.38.60 : SRI-GILLIGAN.ARPA,SRI-GILLIGAN,GILLIGAN,SRI-SAC0 : SUN-120 : UNIX :
TCP/TELNET,TCP/FTP,TCP/SMTP,UDP,ICMP :
HOST : 192.5.89.61 : ZOTZ.ARPA,ZOTZ,SOATS : SUN-2/120 : UNIX : TCP/TELNET,TCP/SMTP,TCP/FTP,ICMP,UDP :
HOST : 192.5.89.62 : TZEC.ARPA,TZEC,ZAKE : SUN-2/120 : UNIX : TCP/TELNET,TCP/SMTP,TCP/FTP,ICMP,UDP :
HOST : 192.5.89.63 : XUL.ARPA,XUL,SHOOL : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.12.5.63 : SRI-MARVIN.ARPA,SRI-MARVIN,MARVIN : SUN-2/170 : UNIX : TCP/SMTP,TCP/TELNET,TCP/FTP :
HOST : 192.5.89.64 : YAXKIN.ARPA,YAXKIN,YOSHKEN : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.65 : MOL.ARPA,MOL,MOLE : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.1.2.66, 8.4.0.12 : BBNF.ARPA,BBNF,BBN-TENEXF : DEC-2040T : TOPS20 :
TCP/FTP,TCP/TELNET,TCP/SMTP,TCP/TIME,UDP/TIME :
HOST : 192.5.89.66 : KANKIN.ARPA,KANKIN,CONKEEN : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.67 : ZIP.ARPA,ZIP,SEEP : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.68 : POP.ARPA,POP,POPE : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.69 : MUAN.ARPA,MUAN,MWAHN : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.70 : UO.ARPA,UO,WOE : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.71 : BUN.ARPA,BUN : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.72 : FUN.ARPA,FUN : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.73 : NUN.ARPA,NUN : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
HOST : 192.5.89.74 : RUN.ARPA,RUN : SUN-2/120 : UNIX : TCP/TELNET,TCP/FTP,TCP/SMTP,ICMP,UDP :
```

DNS: Origins

HOSTS.TXT (TCP/IP, DOMAIN NAME(The Initial and Temporary Domain))

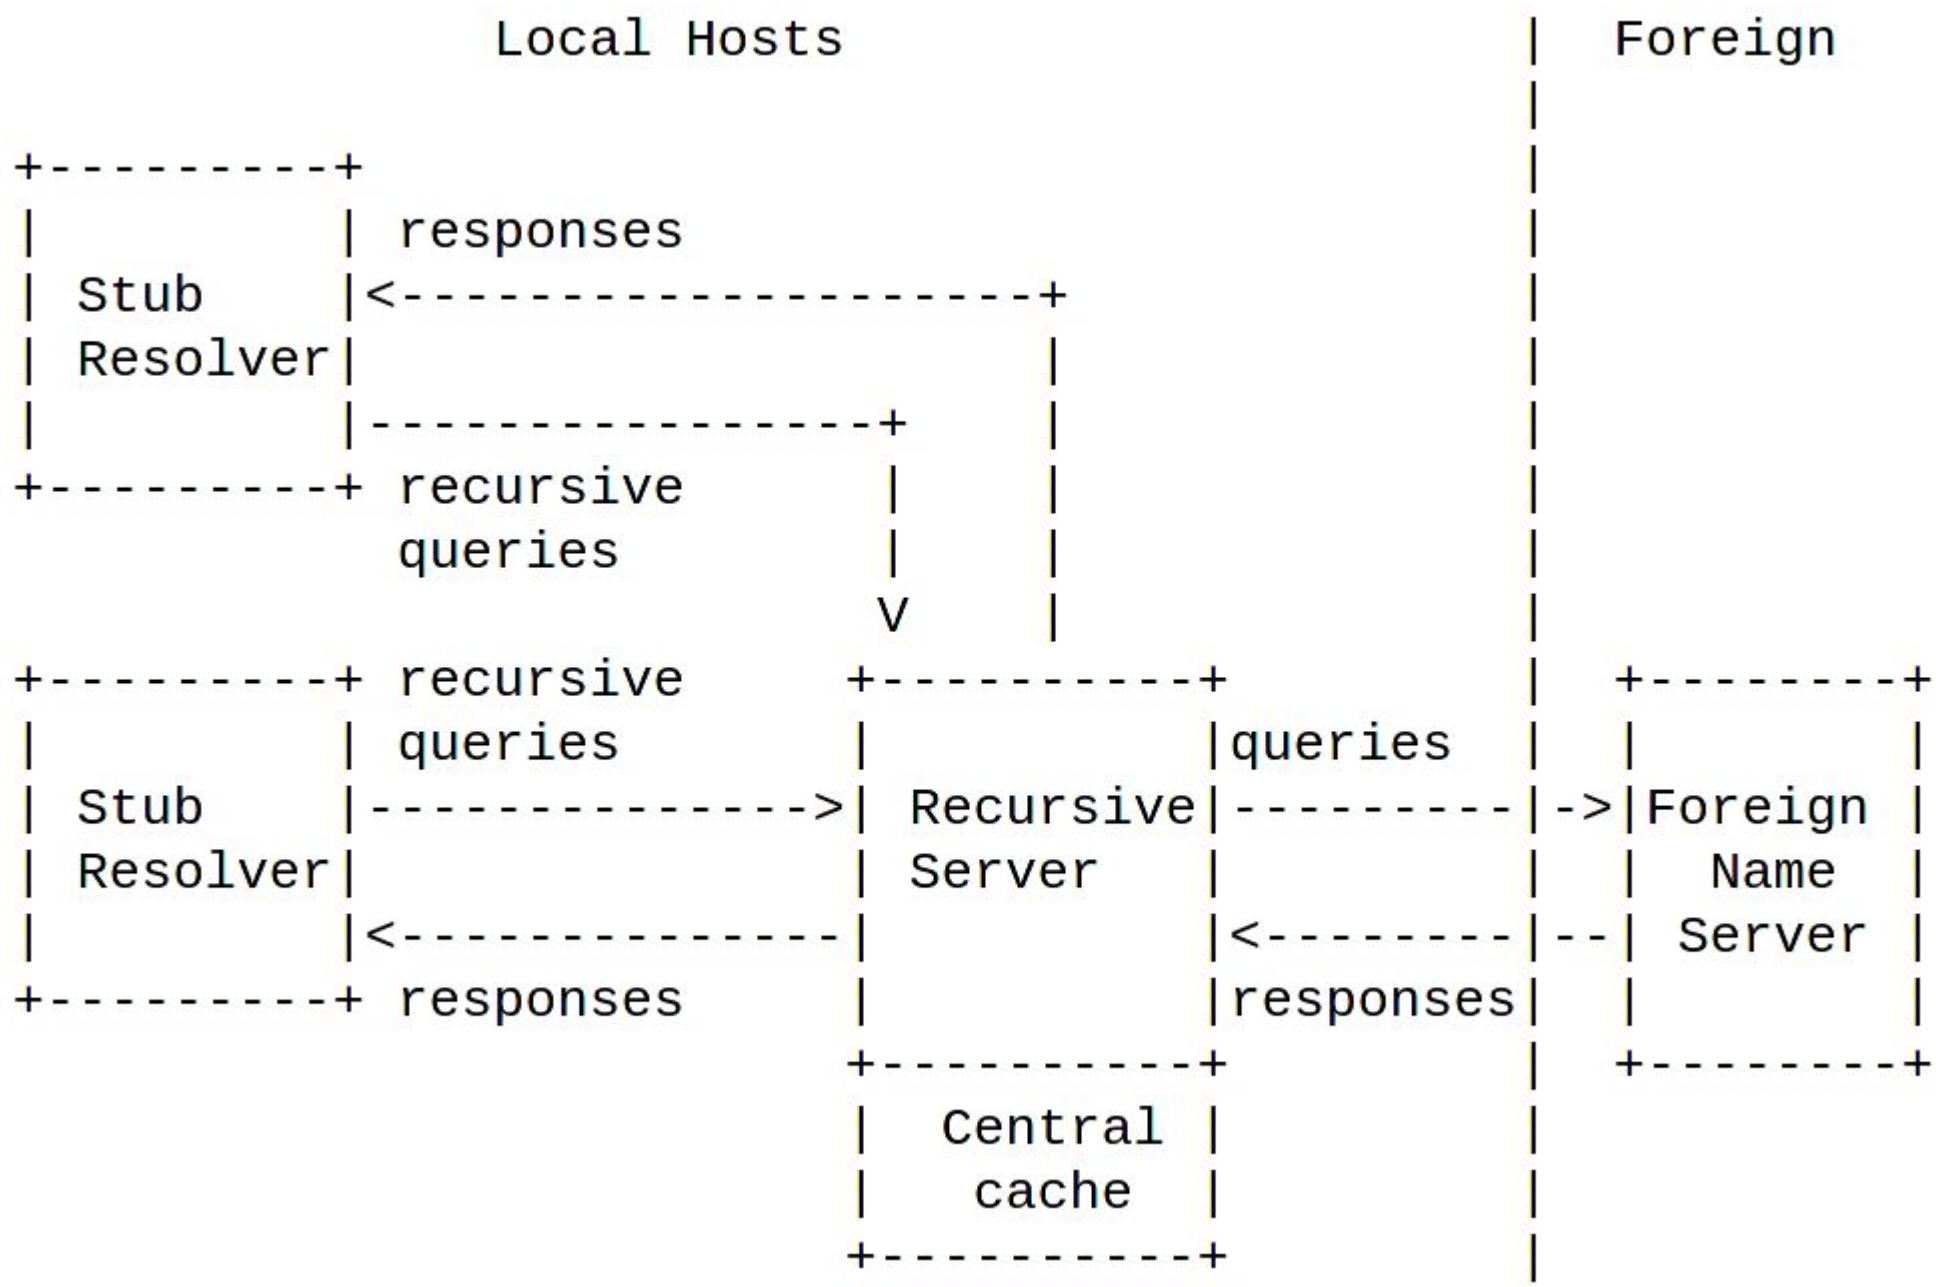
See also RFC 897, RFC 881, and RFC 921.

| Date | Version | File | NET | GATEWAY | HOST |
|------------|---------|--------------------------------|-----|---------|------|
| 1984-03-27 | 340 | HOSTS.TXT (*1) | 135 | 69 | 732 |
| 1984-04-19 | 343 | HOSTS.TXT (*1) | 139 | 72 | 756 |
| 1984-04-27 | 344 | HOSTS.TXT (*1) | 140 | 73 | 766 |
| 1984-06-25 | 353 | HOSTS.TXT (*1) | 144 | 79 | 841 |
| 1984-07-13 | 359 | HOSTS.TXT (*1) | 147 | 82 | 877 |
| 1984-11-28 | 405 | HOSTS.TXT (*1) | 185 | 96 | 1122 |
| 1984-12-21 | 412 | HOSTS.TXT (*1) | 187 | 98 | 1138 |
| 1985-01-02 | 413 | HOSTS.TXT (*1) | 187 | 98 | 1142 |

<https://github.com/ttkzw/hosts.txt>


```
$ cat /etc/hosts
127.0.0.1    localhost loopback
::1         localhost
```

DNS: Origins



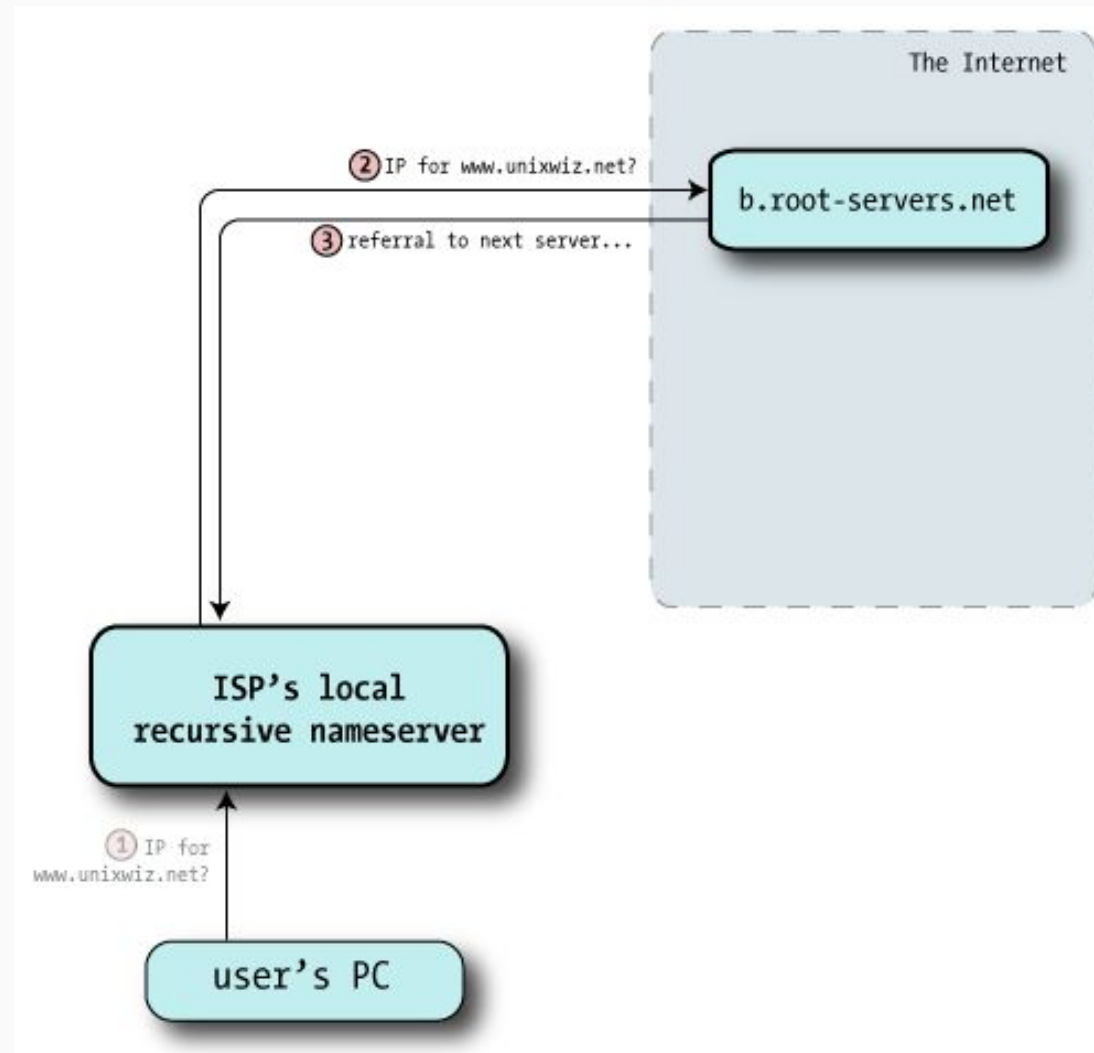
Exercises

Q1.2: DNS Recursive Resolution



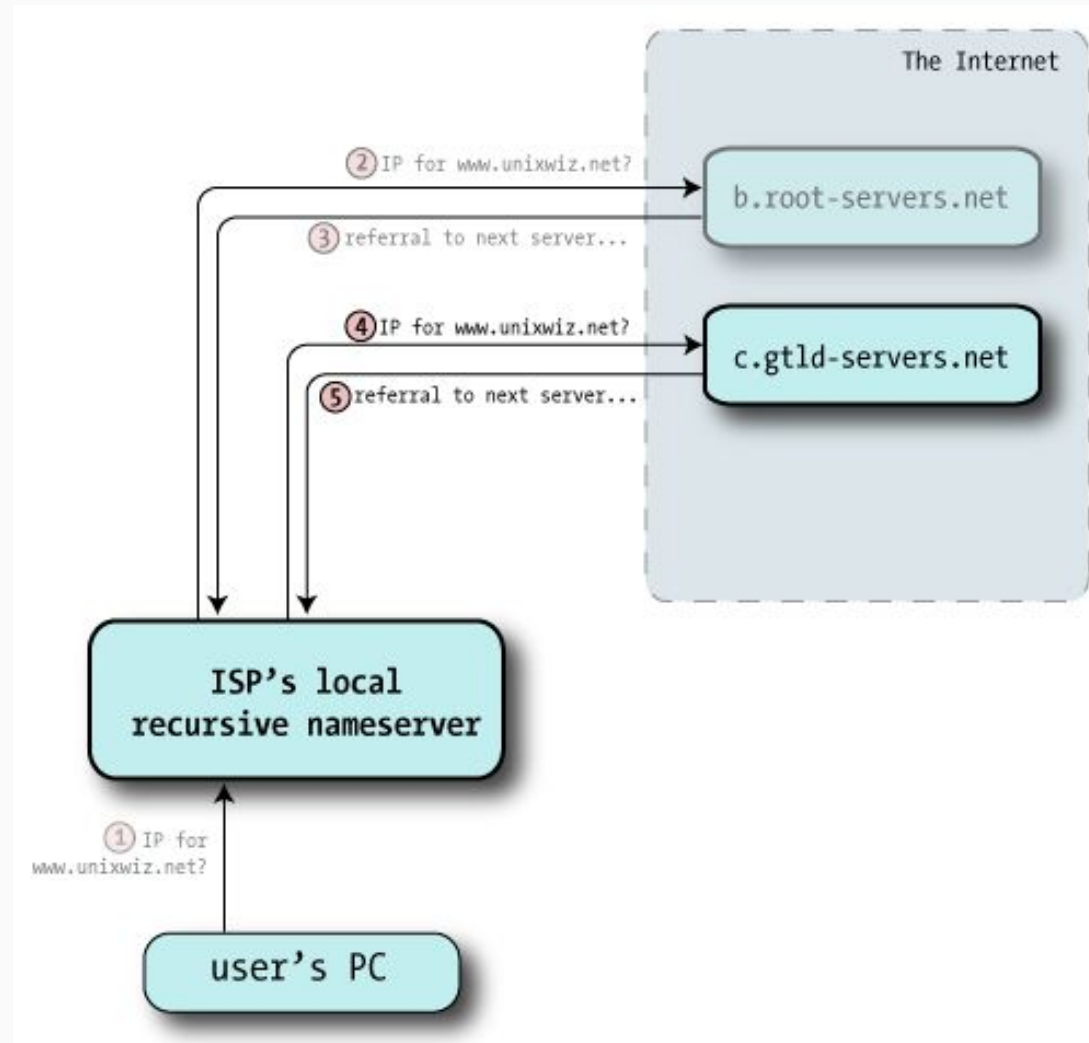
<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

Q1.2: DNS Recursive Resolution



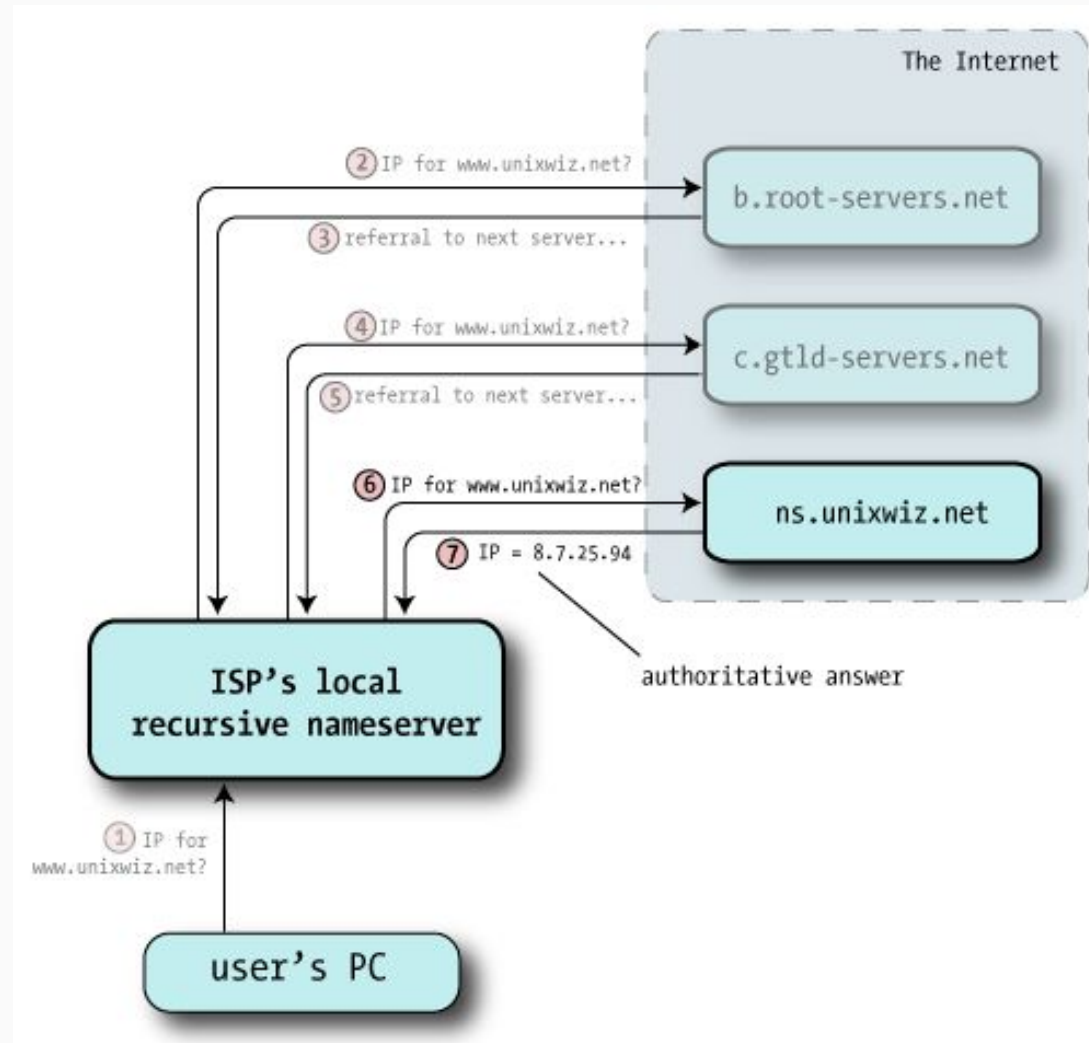
<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

Q1.2: DNS Recursive Resolution




<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

Q1.2: DNS Recursive Resolution



<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

Q1.2: DNS Recursive Resolution



<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

DNS: Glue

```
dig @a.nic.ch. ethz.ch
```

```
;; AUTHORITY SECTION:
```

| | | | | |
|----------|------|----|----|----------------|
| ethz.ch. | 3600 | IN | NS | ns2.ethz.ch. |
| ethz.ch. | 3600 | IN | NS | ns2.switch.ch. |
| ethz.ch. | 3600 | IN | NS | ns1.ethz.ch. |

DNS: Glue

```
dig @a.nic.ch. ethz.ch
```

```
;; AUTHORITY SECTION:
```

| | | | | |
|----------|------|----|----|---------------------|
| ethz.ch. | 3600 | IN | NS | ns2.ethz.ch. |
| ethz.ch. | 3600 | IN | NS | ns2.switch.ch. |
| ethz.ch. | 3600 | IN | NS | ns1.ethz.ch. |

```
;; ADDITIONAL SECTION:
```

| | | | | |
|---------------------|-------------|-----------|-------------|-------------------------|
| ns2.ethz.ch. | 3600 | IN | AAAA | 2001:67c:10ec::b |
| ns2.switch.ch. | 3600 | IN | AAAA | 2001:620:0:ff::2f |
| ns1.ethz.ch. | 3600 | IN | AAAA | 2001:67c:10ec::a |

DNS: Glue

```
dig @a.nic.ch. ethz.ch
```

```
;; AUTHORITY SECTION:
```

| | | | | |
|----------|------|----|----|-----------------------|
| ethz.ch. | 3600 | IN | NS | ns2.ethz.ch. |
| ethz.ch. | 3600 | IN | NS | ns2.switch.ch. |
| ethz.ch. | 3600 | IN | NS | ns1.ethz.ch. |

```
;; ADDITIONAL SECTION:
```

| | | | | |
|-----------------------|-------------|-----------|-------------|--------------------------|
| ns2.ethz.ch. | 3600 | IN | AAAA | 2001:67c:10ec::b |
| ns2.switch.ch. | 3600 | IN | AAAA | 2001:620:0:ff::2f |
| ns1.ethz.ch. | 3600 | IN | AAAA | 2001:67c:10ec::a |

DNS: Glue

```
dig @ns1.attacker.com. attacker.com
```

```
;; AUTHORITY SECTION:
```

| | | | | |
|---------------|------|----|----|-----------------------|
| attacker.com. | 3600 | IN | NS | ns2.attacker.com |
| attacker.com. | 3600 | IN | NS | ns1.google.com |

```
;; ADDITIONAL SECTION:
```

| | | | | |
|-----------------------|-------------|-----------|----------|----------------|
| ns2.attacker.com | 3600 | IN | A | 1.3.3.7 |
| ns1.google.com | 3600 | IN | A | 1.3.3.7 |

Q2.1: DNS Lame Delegations

"authoritative" response. If not, this is called a "lame delegation". A lame delegation exists when a nameserver is delegated responsibility for providing nameservice for a zone (via NS records) but is not performing nameservice for that zone (usually because it is not set up as a primary or secondary for the zone).

The "classic" lame delegation can be illustrated in this example:

| | | | |
|------------|----|----|-----------------|
| podunk.xx. | IN | NS | ns1.podunk.xx. |
| | IN | NS | ns0.widget.com. |

"podunk.xx" is a new domain which has recently been created, and "ns1.podunk.xx" has been set up to perform nameservice for the zone. They haven't quite finished everything yet and haven't made sure that the hostmaster at "ns0.widget.com" has set up to be a proper secondary, and thus has no information about the podunk.xx domain.

<https://www.ietf.org/rfc/rfc1912.txt>

DNS: Glue

```
dig @a.nic.ch. ethz.ch
```

```
;; AUTHORITY SECTION:
```

| | | | | |
|----------|------|----|----|-----------------------|
| ethz.ch. | 3600 | IN | NS | ns2.ethz.ch. |
| ethz.ch. | 3600 | IN | NS | ns2.switch.ch. |
| ethz.ch. | 3600 | IN | NS | ns1.ethz.ch. |

```
;; ADDITIONAL SECTION:
```

| | | | | |
|-----------------------|-------------|-----------|-------------|--------------------------|
| ns2.ethz.ch. | 3600 | IN | AAAA | 2001:67c:10ec::b |
| ns2.switch.ch. | 3600 | IN | AAAA | 2001:620:0:ff::2f |
| ns1.ethz.ch. | 3600 | IN | AAAA | 2001:67c:10ec::a |

DNS: Glue

```
dig @ns1.attacker.com. attacker.com
;; AUTHORITY SECTION:
attacker.com.          3600    IN      NS      ns2.attacker.com
attacker.com.          3600    IN      NS      ns1.google.com

;; ADDITIONAL SECTION:
ns2.attacker.com        3600    IN      A       1.3.3.7
ns1.google.com          3600    IN      A       1.3.3.7
```

Q2.1: DNS Lame Delegations

"authoritative" response. If not, this is called a "lame delegation". A lame delegation exists when a nameserver is delegated responsibility for providing nameservice for a zone (via NS records) but is not performing nameservice for that zone (usually because it is not set up as a primary or secondary for the zone).

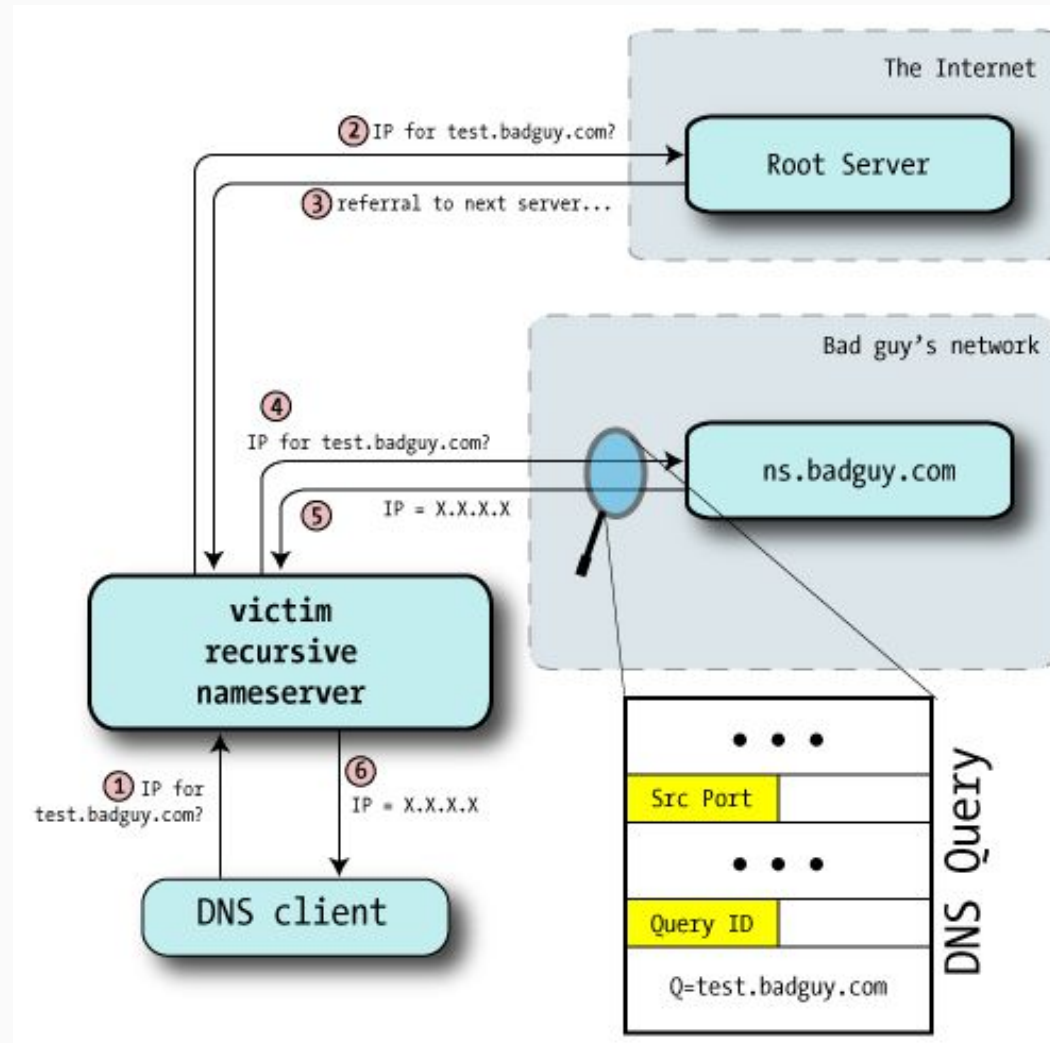
The "classic" lame delegation can be illustrated in this example:

```
podunk.xx.      IN      NS      ns1.podunk.xx.  
                ns0.widget.com.
```

podunk.xx is a new domain which has recently been created, and "ns1.podunk.xx" has been set up to perform nameservice for the zone. They haven't quite finished everything yet and haven't made sure that the hostmaster at "ns0.widget.com" has set up to be a proper secondary, and thus has no information about the podunk.xx domain.

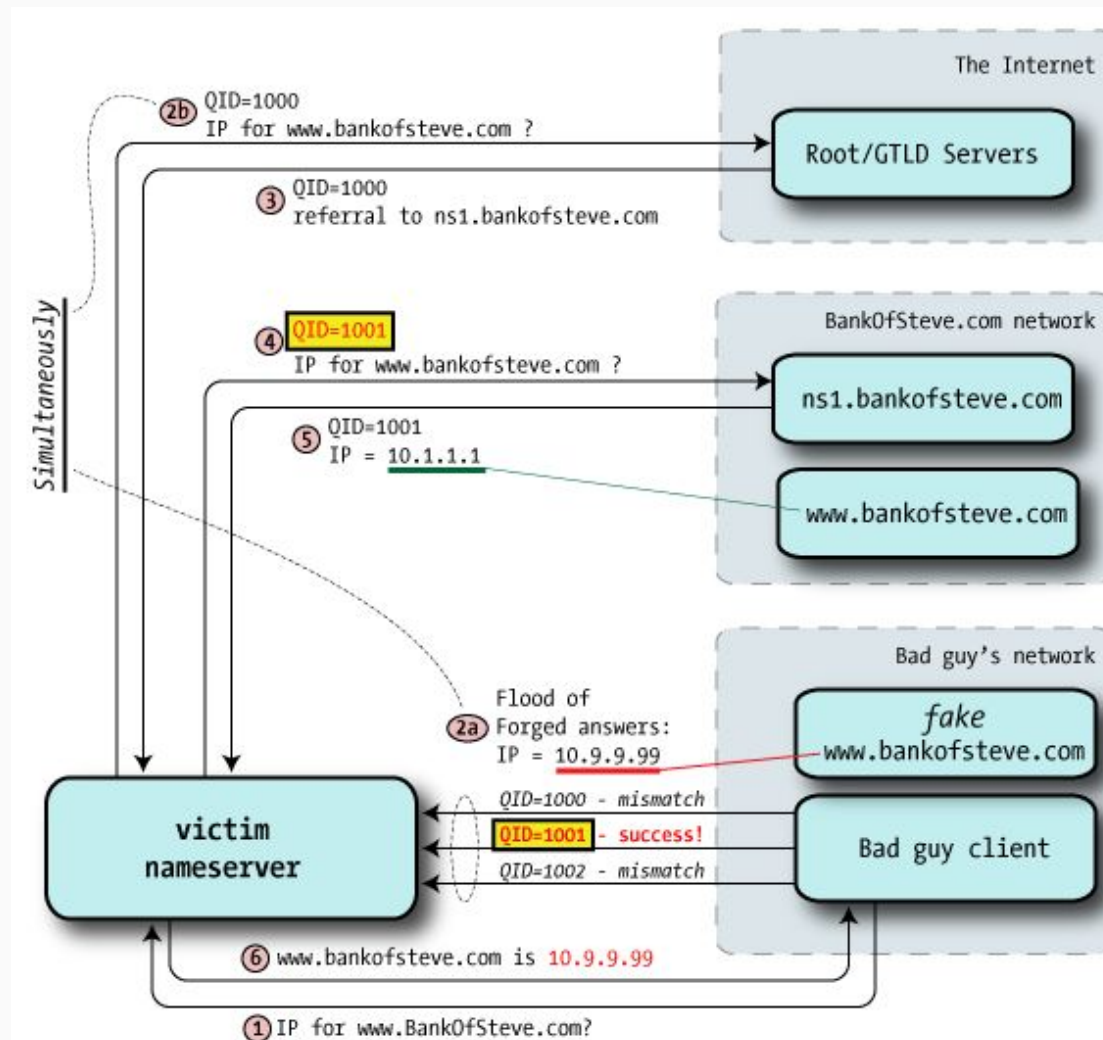
<https://www.ietf.org/rfc/rfc1912.txt>

DNS: 2009 Kaminsky Attack



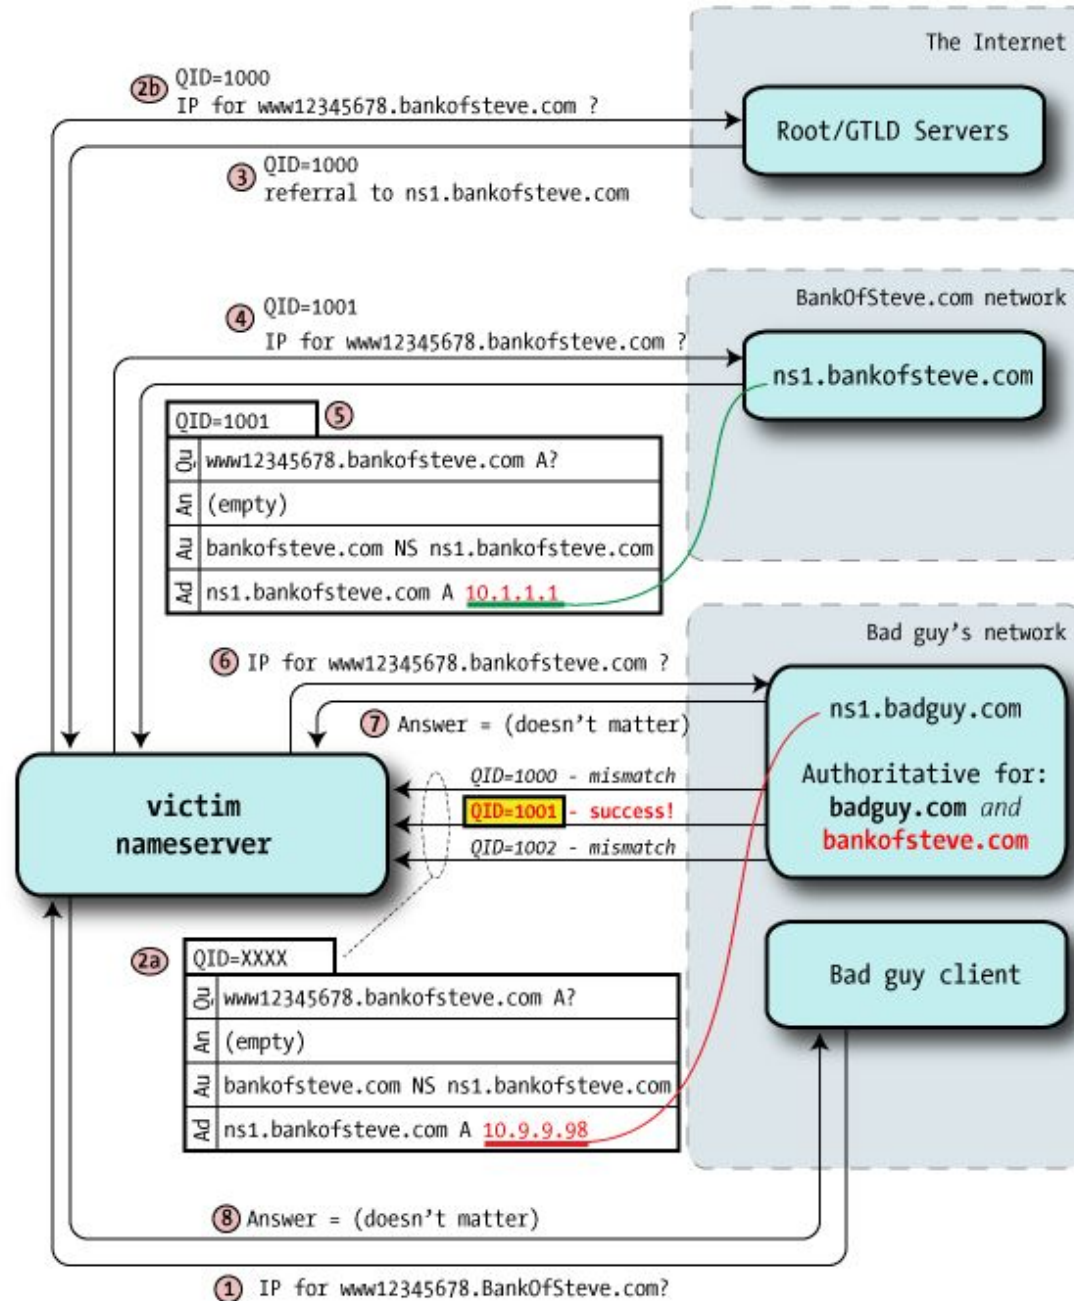
<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

DNS: 2009 Kaminsky Attack



<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

DNS: 2009 Kaminsky Attack



ixwiz.net/techtips/iguide-dns-vuln.html

Use random ports!

$$\underbrace{2^{16}}_{\text{Query ID}} \times \underbrace{2^{11}}_{\text{Source ports}} = 2^{27} = \mathbf{134 \text{ million}}$$

<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

Use random ports

SOLVED ~2009

$2^{16} \times 2^{11} = 2^{27} = 134 \text{ million}$

Source ports
Query ID

<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

DNS: 2009 Fix

$$\underbrace{2^{16}}_{\text{Query ID}} \times \underbrace{2^{11}}_{\text{Source ports}} = 2^{27} = \mathbf{134 \text{ million}}$$

DJB Was Right

One nameserver is notable for having gotten **both** the query-id and source-port randomness right from the start: [DJBDNS](#) by the legendary Daniel J. Bernstein.

Though long a lightning rod for controversy, he's clearly walked the walk on security: there's been just [one minor security vulnerability](#) in DJBDNS.


<http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>

Use random ports!


$$\underbrace{2^{16}}_{\text{Query ID}} \times \underbrace{2^{11}}_{\text{Source ports}} = 2^{27} = \mathbf{134 \text{ million}}$$

Q2.2.2: Fragmentation


IP Fragmentation



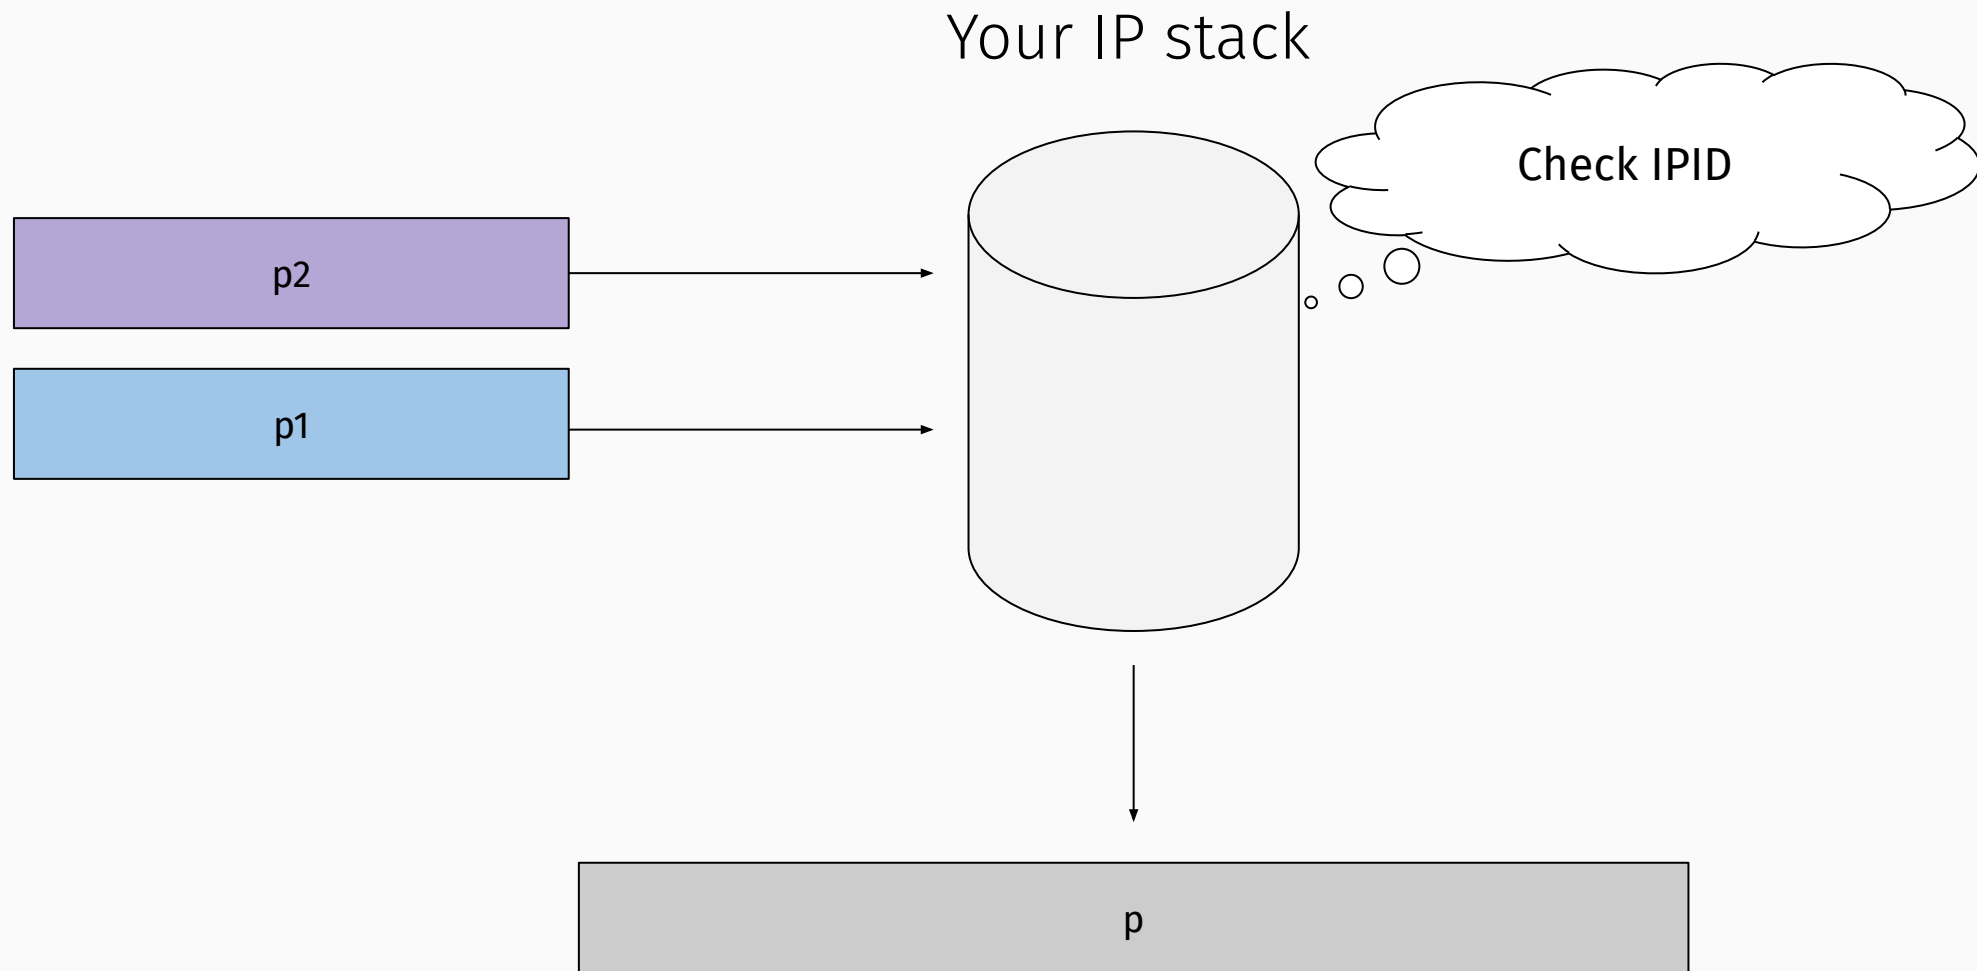
IP Fragmentation



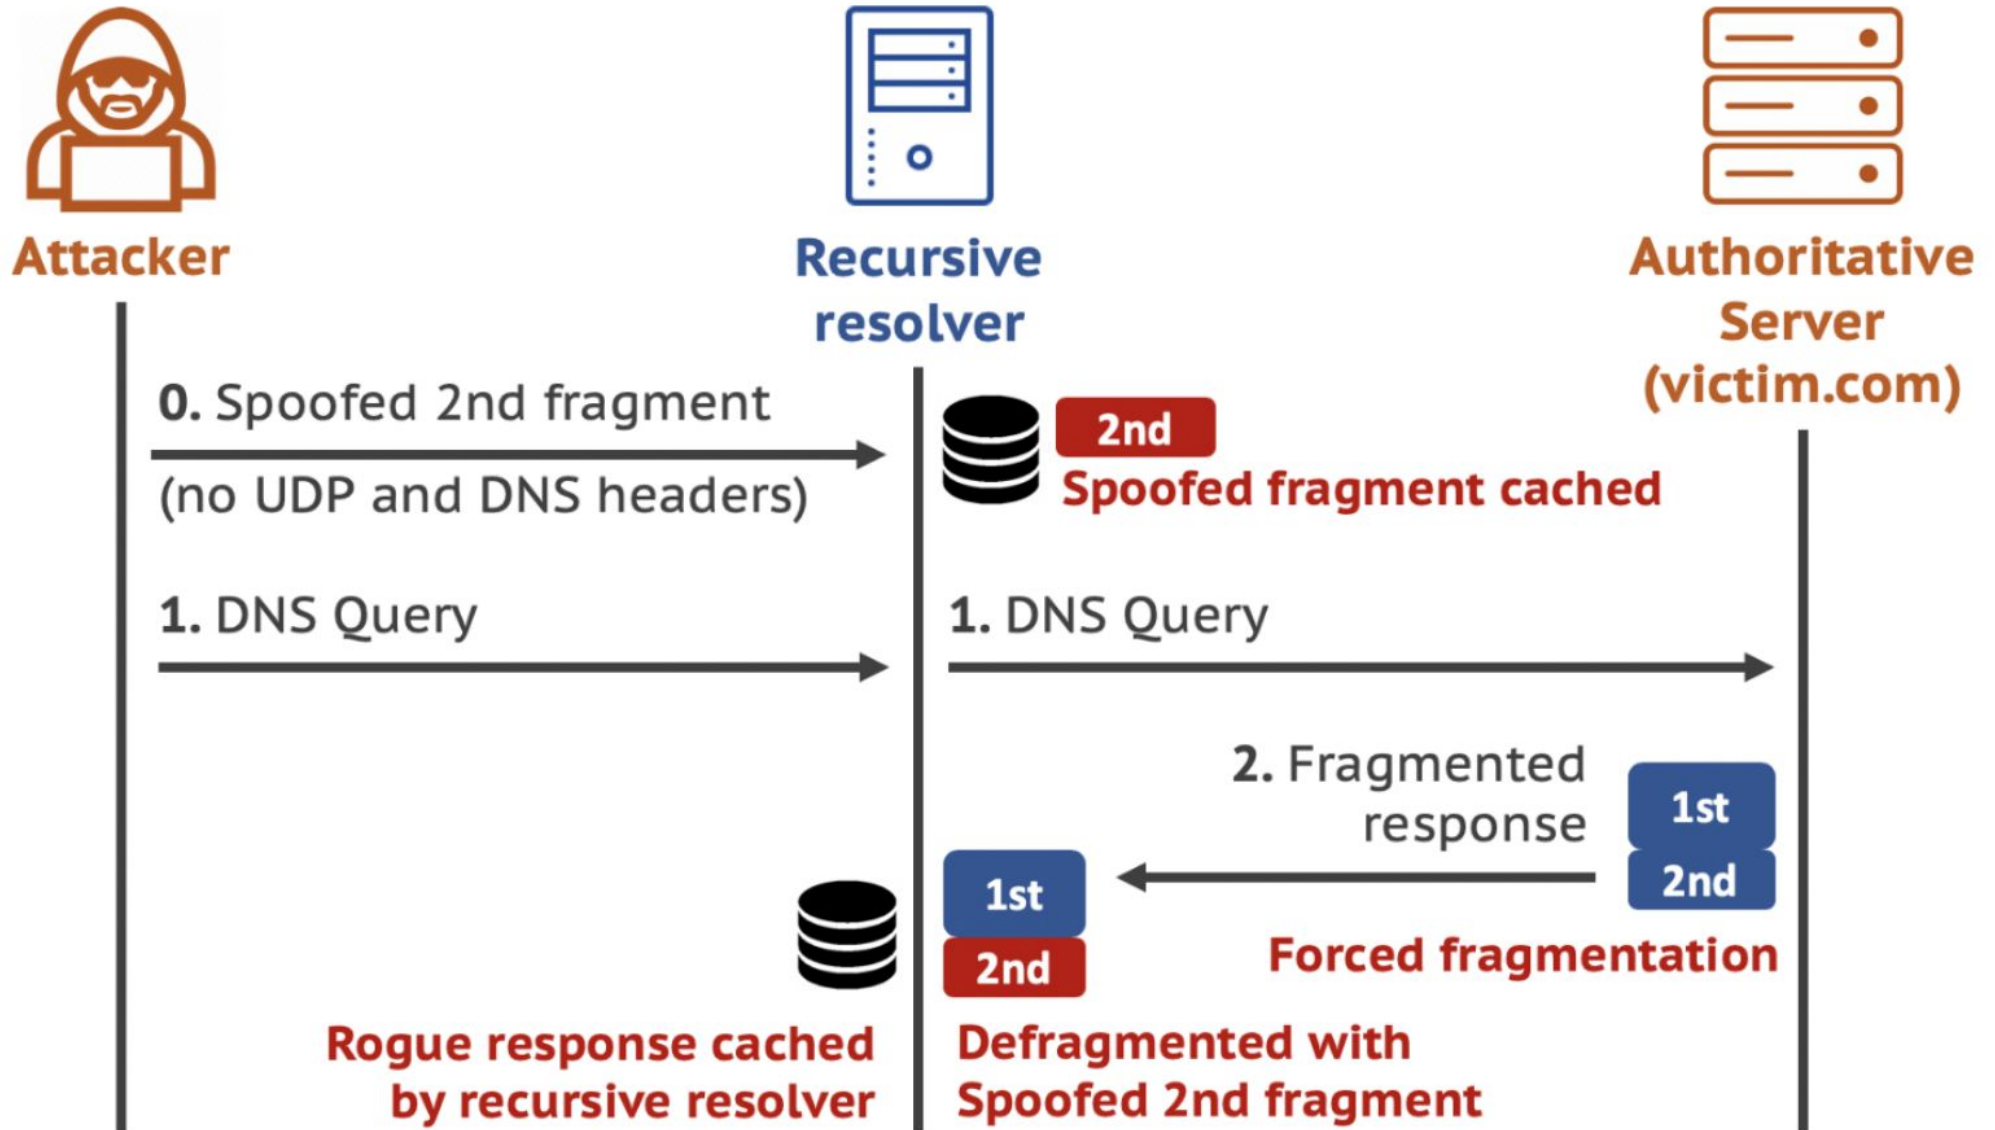
IP Fragmentation



IP Fragmentation



Q2.2.2: Fragmentation



Q2.2.2: Fragmentation

$$\underbrace{2^{16}}_{\text{Query ID}} \times \underbrace{2^{11}}_{\text{Source ports}} = 2^{27} = \mathbf{134 \text{ million}}$$

Q2.2.2: Fragmentation

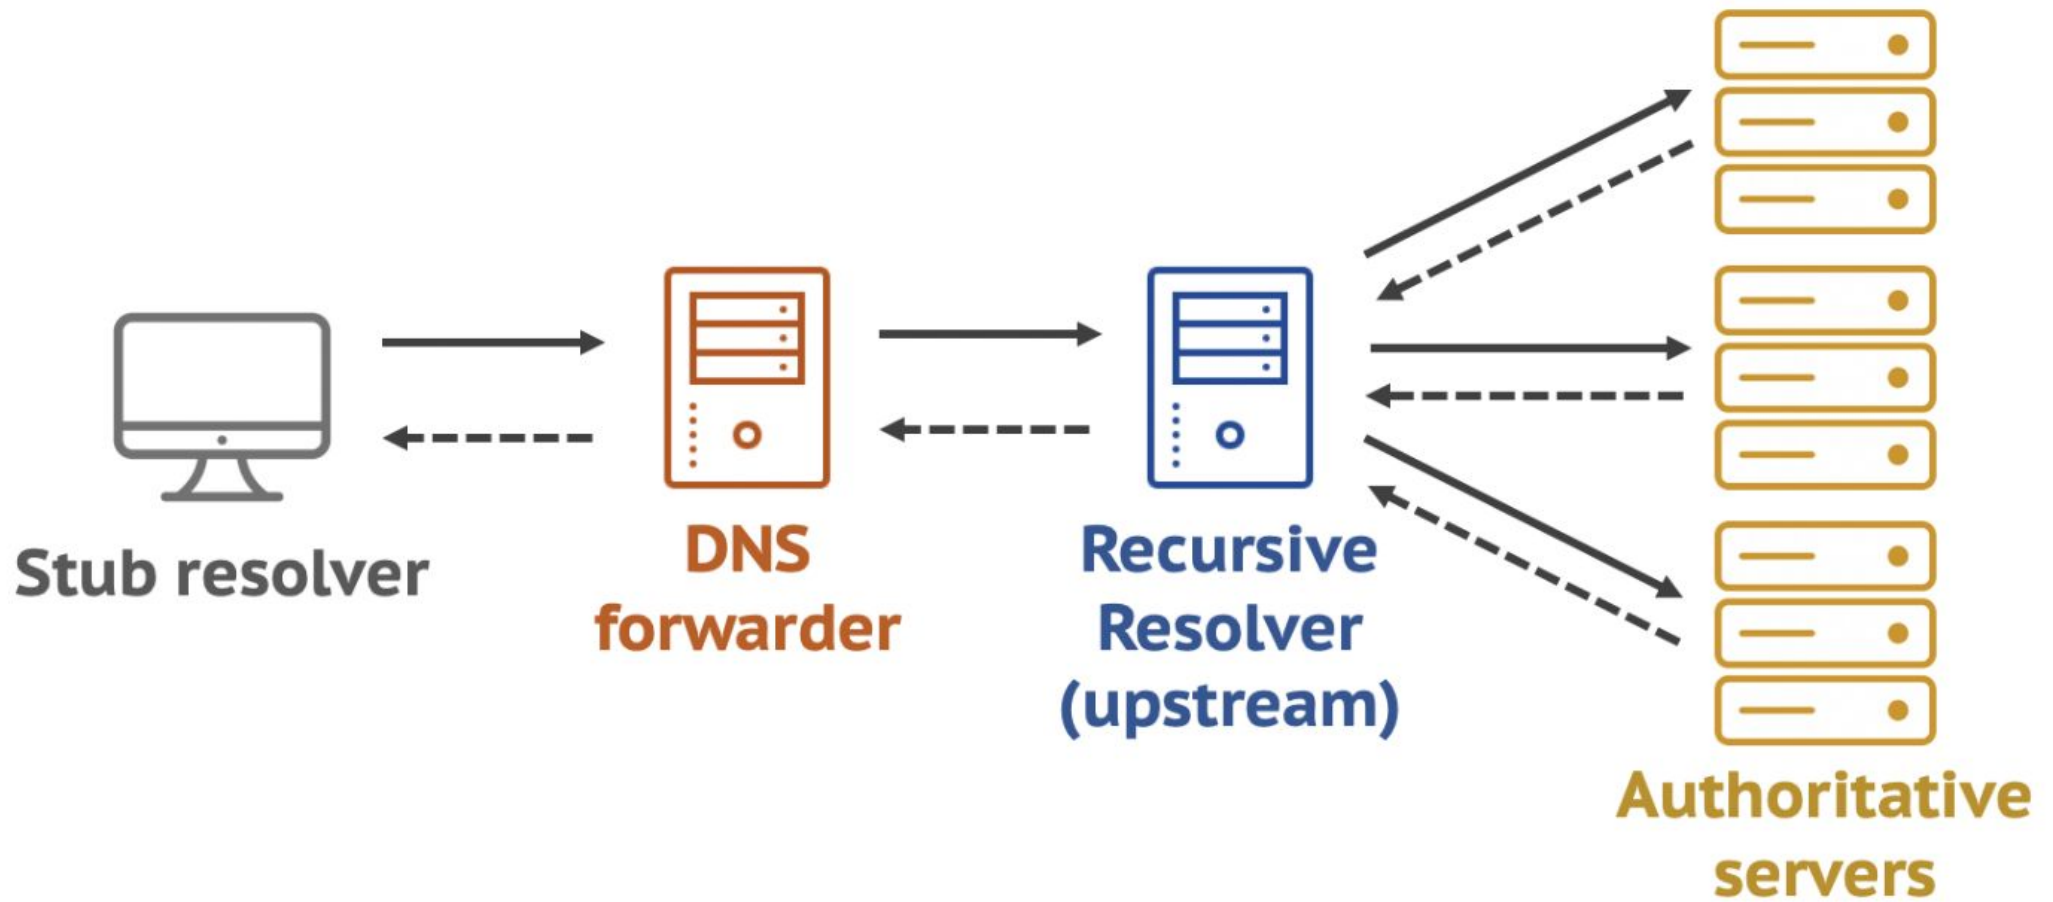
$$\overline{2^{16}} \times \overline{2^{11}} = \overline{2^{27}} \quad 134 \text{ million}$$

Source ports
Query ID

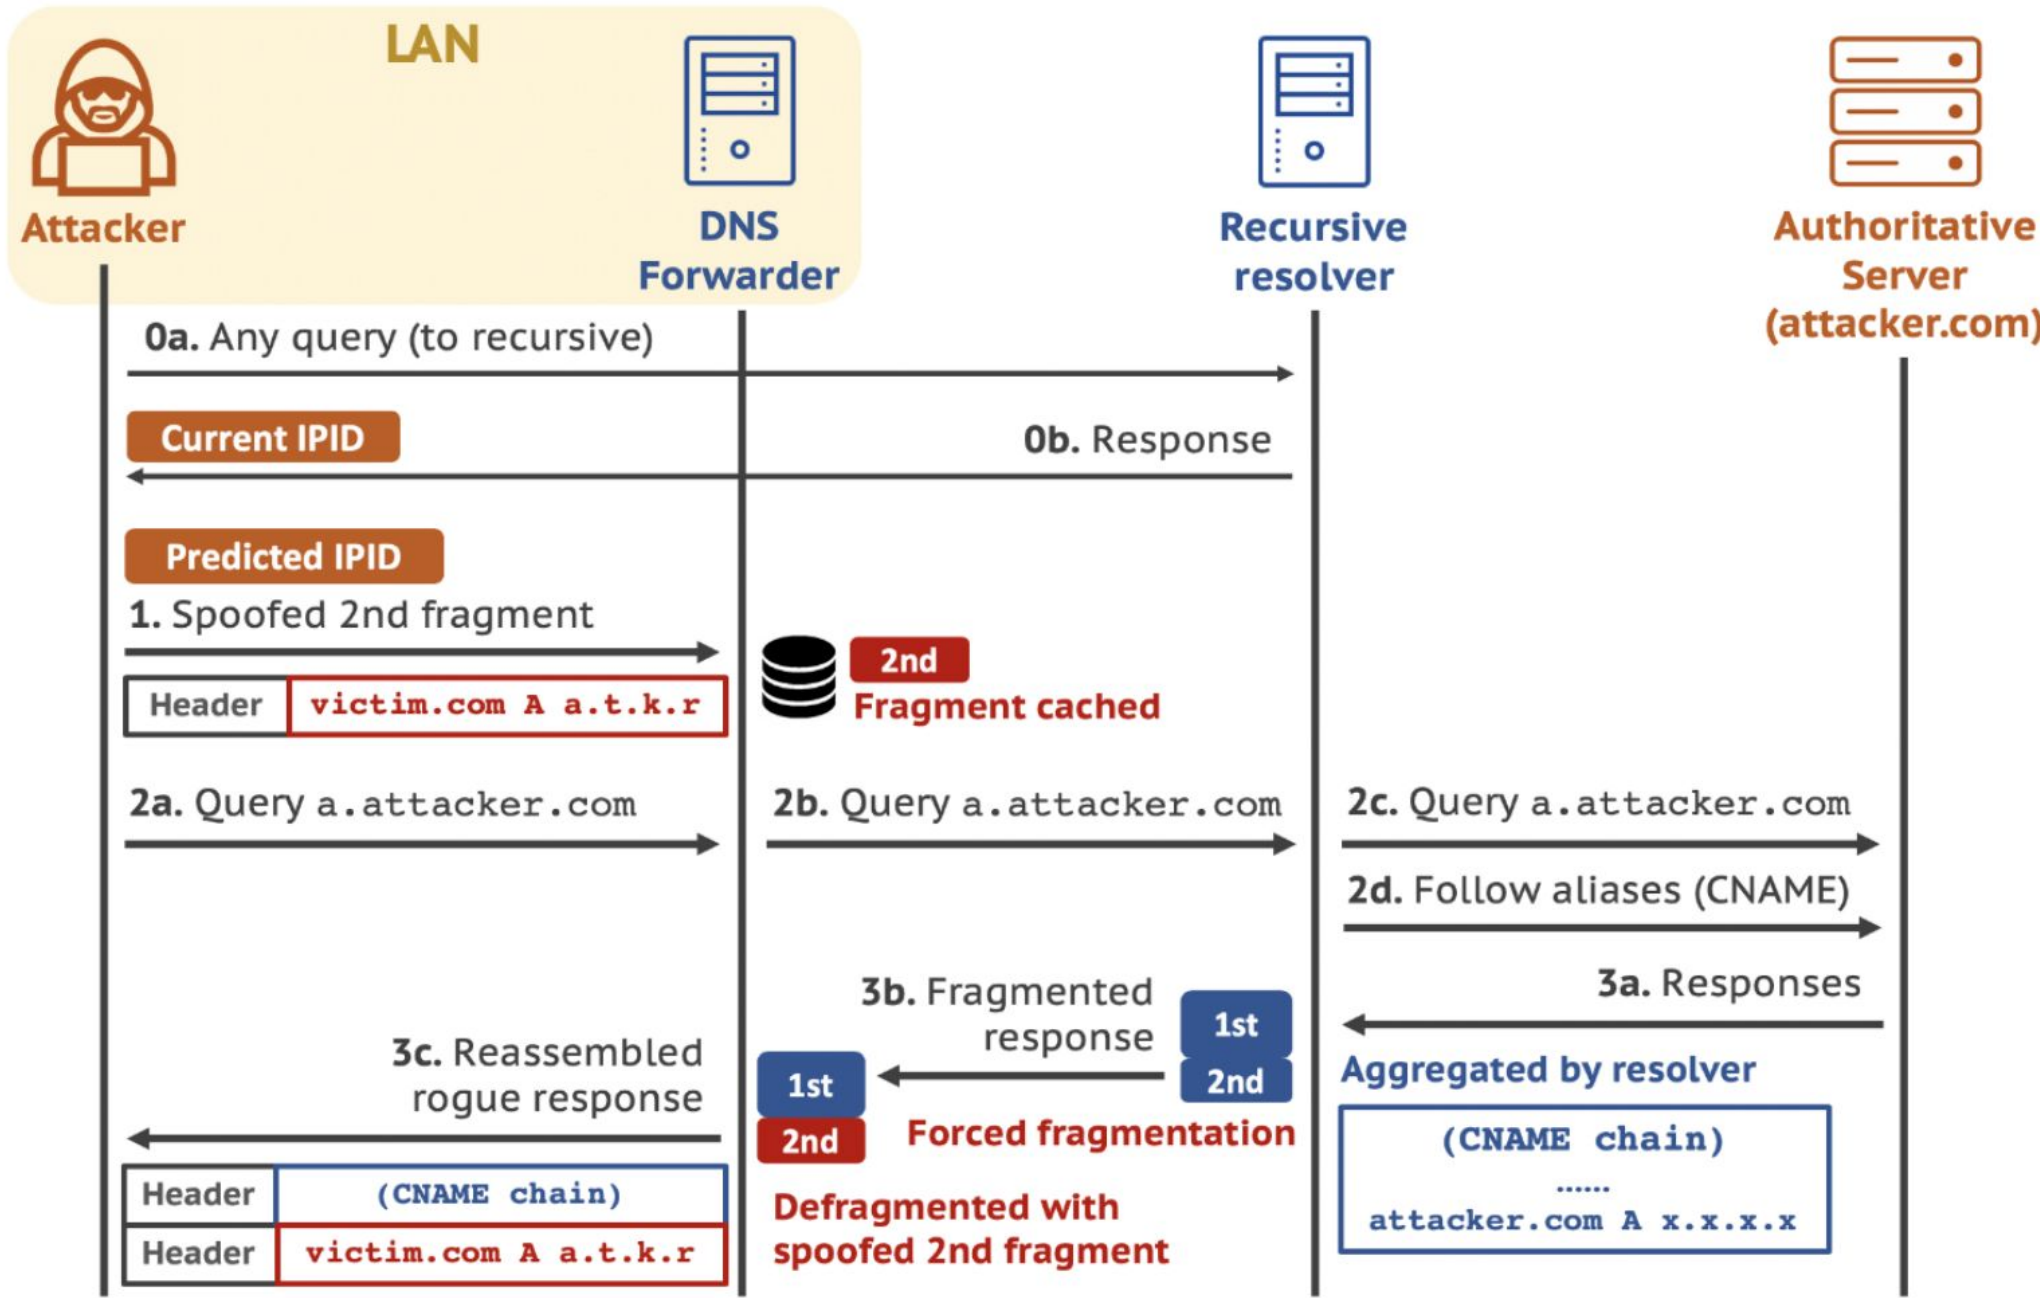
+ IPID (predictable)

Q2.3: Forwarders

Q2.3: Forwarders



Q2.4: Fragmenting Forwarder



Q2.4: Fragmenting Forwarder

1st fragment

| |
|-------------------------------------|
| a.attacker.com CNAME b.attacker.com |
| b.attacker.com CNAME c.attacker.com |
| c.attacker.com CNAME d.attacker.com |
| ... |

| |
|-------------------------------------|
| x.attacker.com CNAME y.attacker.com |
| y.attacker.com CNAME z.attacker.com |
| z.attacker.com A x.x.x.x |

2nd fragment

(a) Response received by
upstream recursive resolver



1st fragment

| |
|-------------------------------------|
| a.attacker.com CNAME b.attacker.com |
| b.attacker.com CNAME c.attacker.com |
| c.attacker.com CNAME d.attacker.com |
| ... |

| |
|-------------------------------------|
| x.attacker.com CNAME y.attacker.com |
| y.attacker.com CNAME victim.com |
| victim.com A a.t.k.r |

Spoofed 2nd fragment

(b) Response received by
DNS forwarder

Q2.5: Attacking Recursive Resolvers?

Q2.5: Attacking Recursive Resolvers?

| |
|-------------------------------------|
| x.attacker.com CNAME y.attacker.com |
| y.attacker.com CNAME victim.com |
| victim.com A a.t.k.r |

Spoofed 2nd fragment

**(b) Response received by
DNS forwarder**

"authoritative" response. If not, this is called a "lame delegation". A lame delegation exists when a nameserver is delegated responsibility for providing nameservice for a zone (via NS records) but is not performing nameservice for that zone (usually because it is not set up as a primary or secondary for the zone).

The "classic" lame delegation can be illustrated in this example:

| | | | |
|------------|----|----|-----------------|
| podunk.xx. | IN | NS | ns1.podunk.xx. |
| | IN | NS | ns0.widget.com. |

"podunk.xx" is a new domain which has recently been created, and "ns1.podunk.xx" has been set up to perform nameservice for the zone. They haven't quite finished everything yet and haven't made sure that the hostmaster at "ns0.widget.com" has set up to be a proper secondary, and thus has no information about the podunk.xx domain.

Q2.5: Attacking Recursive Resolvers?

| |
|--|
| <code>x.attacker.com CNAME y.attacker.com</code> |
| <code>y.attacker.com CNAME victim.com</code> |
| <code>victim.com A a.t.k.r</code> |

Spoofted 2nd fragment

**(b) Response received by
DNS forwarder**

"authoritative" DNS server. If not, this is called a "lame delegation". A lame delegation exists when a nameserver is delegated responsibility for providing nameservice for a zone (via NS records) but is not performing nameservice for that zone (usually because it is not set up as a primary or secondary for the zone).

The "classic" lame delegation can be illustrated in this example:

| | | | |
|------------|----|----|-----------------|
| podunk.xx. | IN | NS | ns1.podunk.xx. |
| | IN | NS | ns0.widget.com. |

"podunk.xx" is a new domain which has recently been created, and "ns1.podunk.xx" has been set up to perform nameservice for the zone. They haven't quite finished everything yet and haven't made sure that the hostmaster at "ns0.widget.com" has set up to be a proper secondary, and thus has no information about the podunk.xx domain.

Q2.5: Attacking Recursive Resolvers?

1st fragment

| |
|-------------------------------------|
| a.attacker.com CNAME b.attacker.com |
| b.attacker.com CNAME c.attacker.com |
| c.attacker.com CNAME d.attacker.com |
| ... |

| |
|-------------------------------------|
| x.attacker.com CNAME y.attacker.com |
| y.attacker.com CNAME z.attacker.com |
| z.attacker.com A x.x.x.x |

2nd fragment

(a) Response received by
upstream recursive resolver



1st fragment

| |
|-------------------------------------|
| a.attacker.com CNAME b.attacker.com |
| b.attacker.com CNAME c.attacker.com |
| c.attacker.com CNAME d.attacker.com |
| ... |

| |
|-------------------------------------|
| x.attacker.com CNAME y.attacker.com |
| y.attacker.com CNAME victim.com |
| victim.com A a.t.k.r |

Spoofed 2nd fragment

(b) Response received by
DNS forwarder

Q2.6: Mitigations

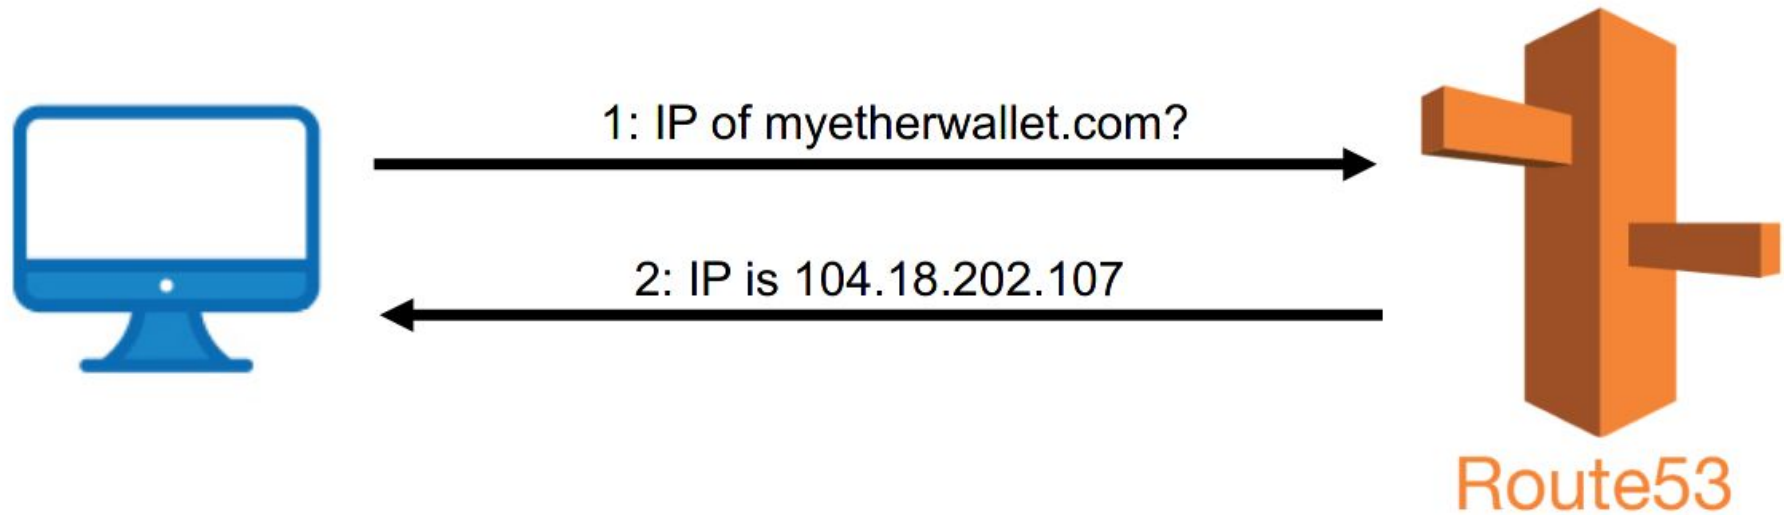
Q2.6: Mitigations

| | | |
|----------------|-------|----------------|
| x.attacker.com | CNAME | y.attacker.com |
| y.attacker.com | CNAME | victim.com |
| victim.com | A | a.t.k.r |

Caching by record: victim.com -> a.t.k.r

Caching by query: victim.com -> new query
attacker.com -> ca

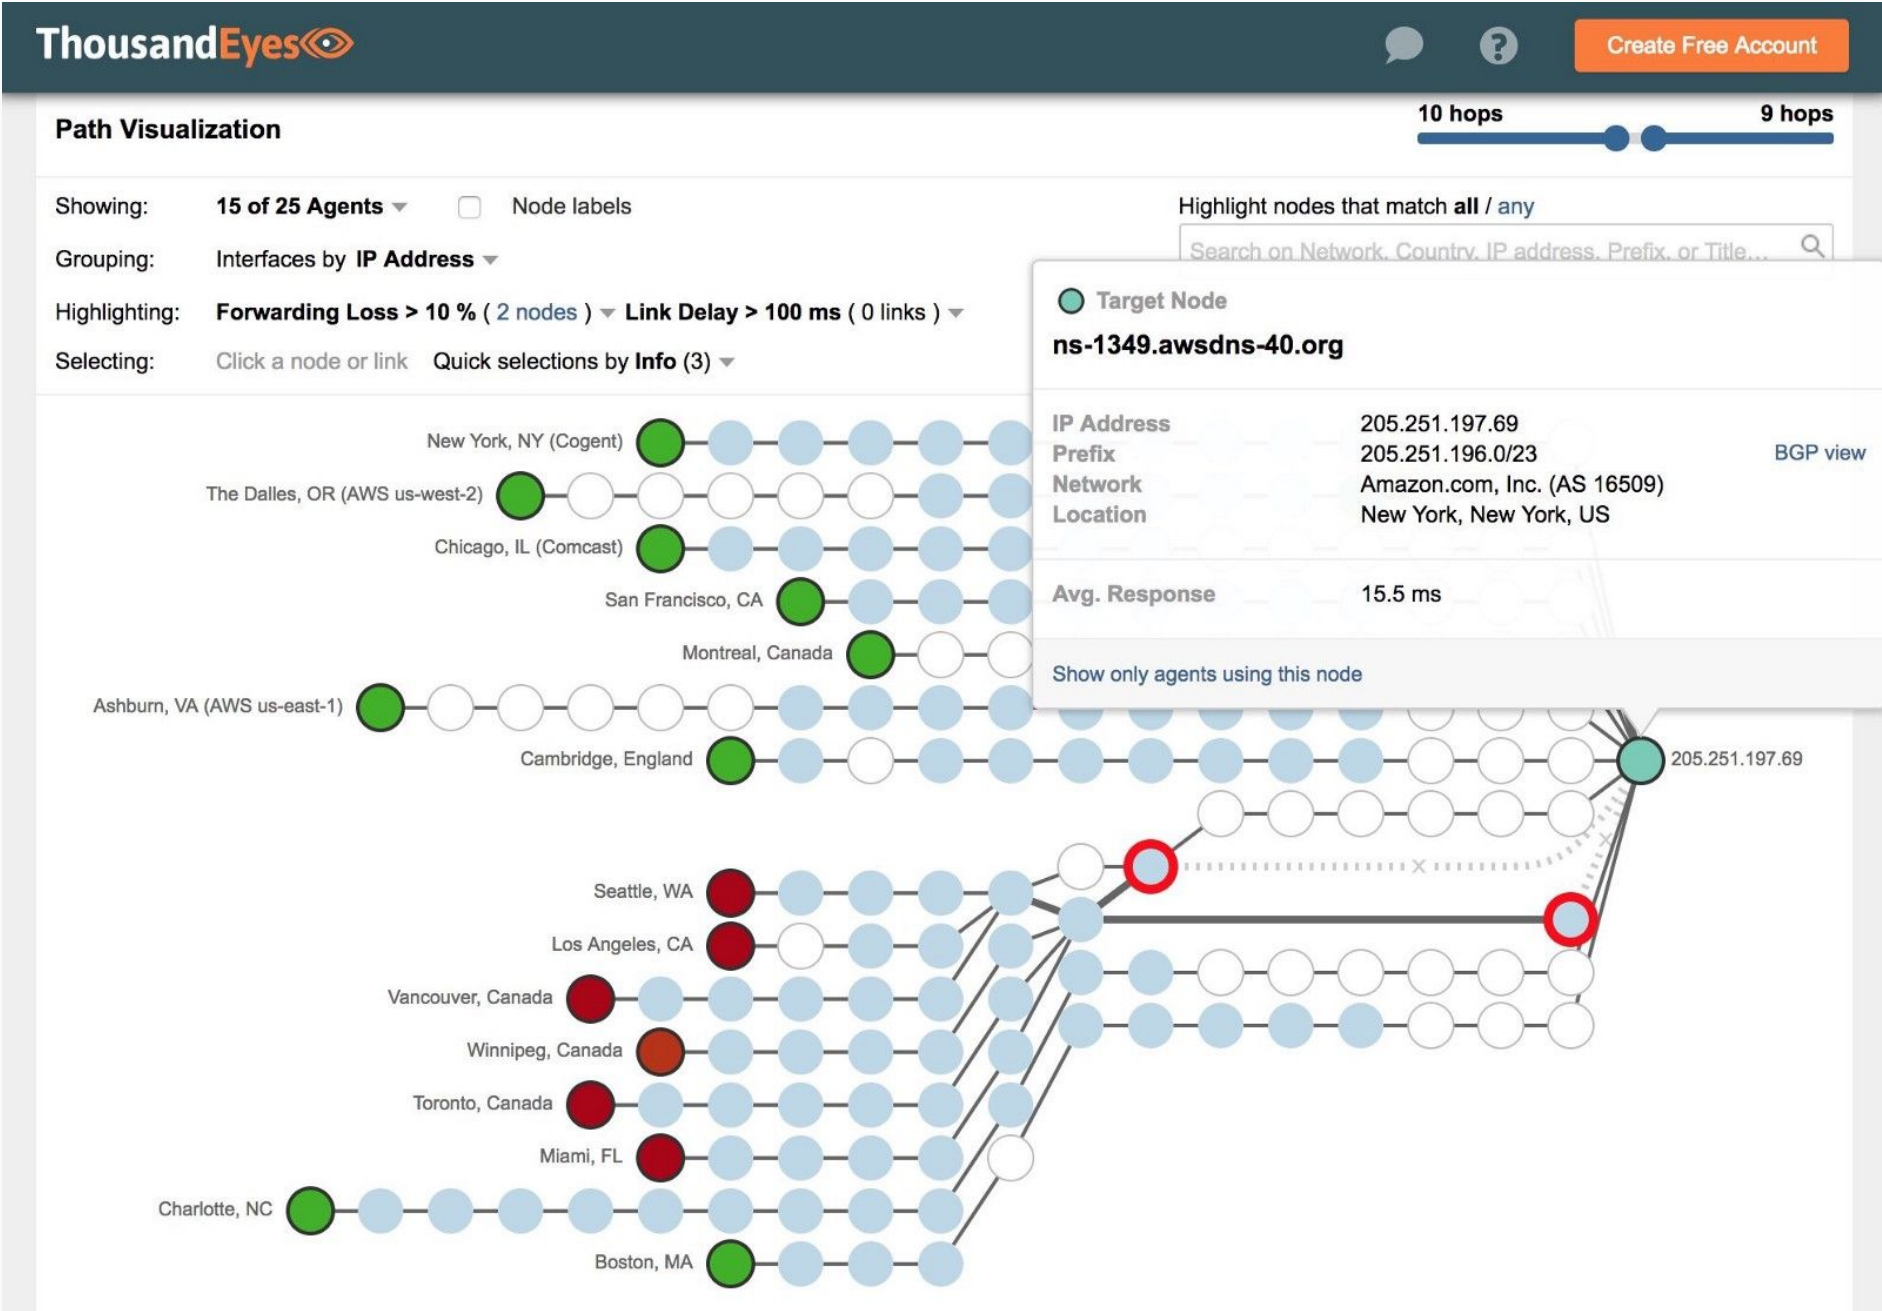
Q3: MyEtherWallet



Authoritative servers of MyEtherWallet hosted at
Amazon AWS Route53

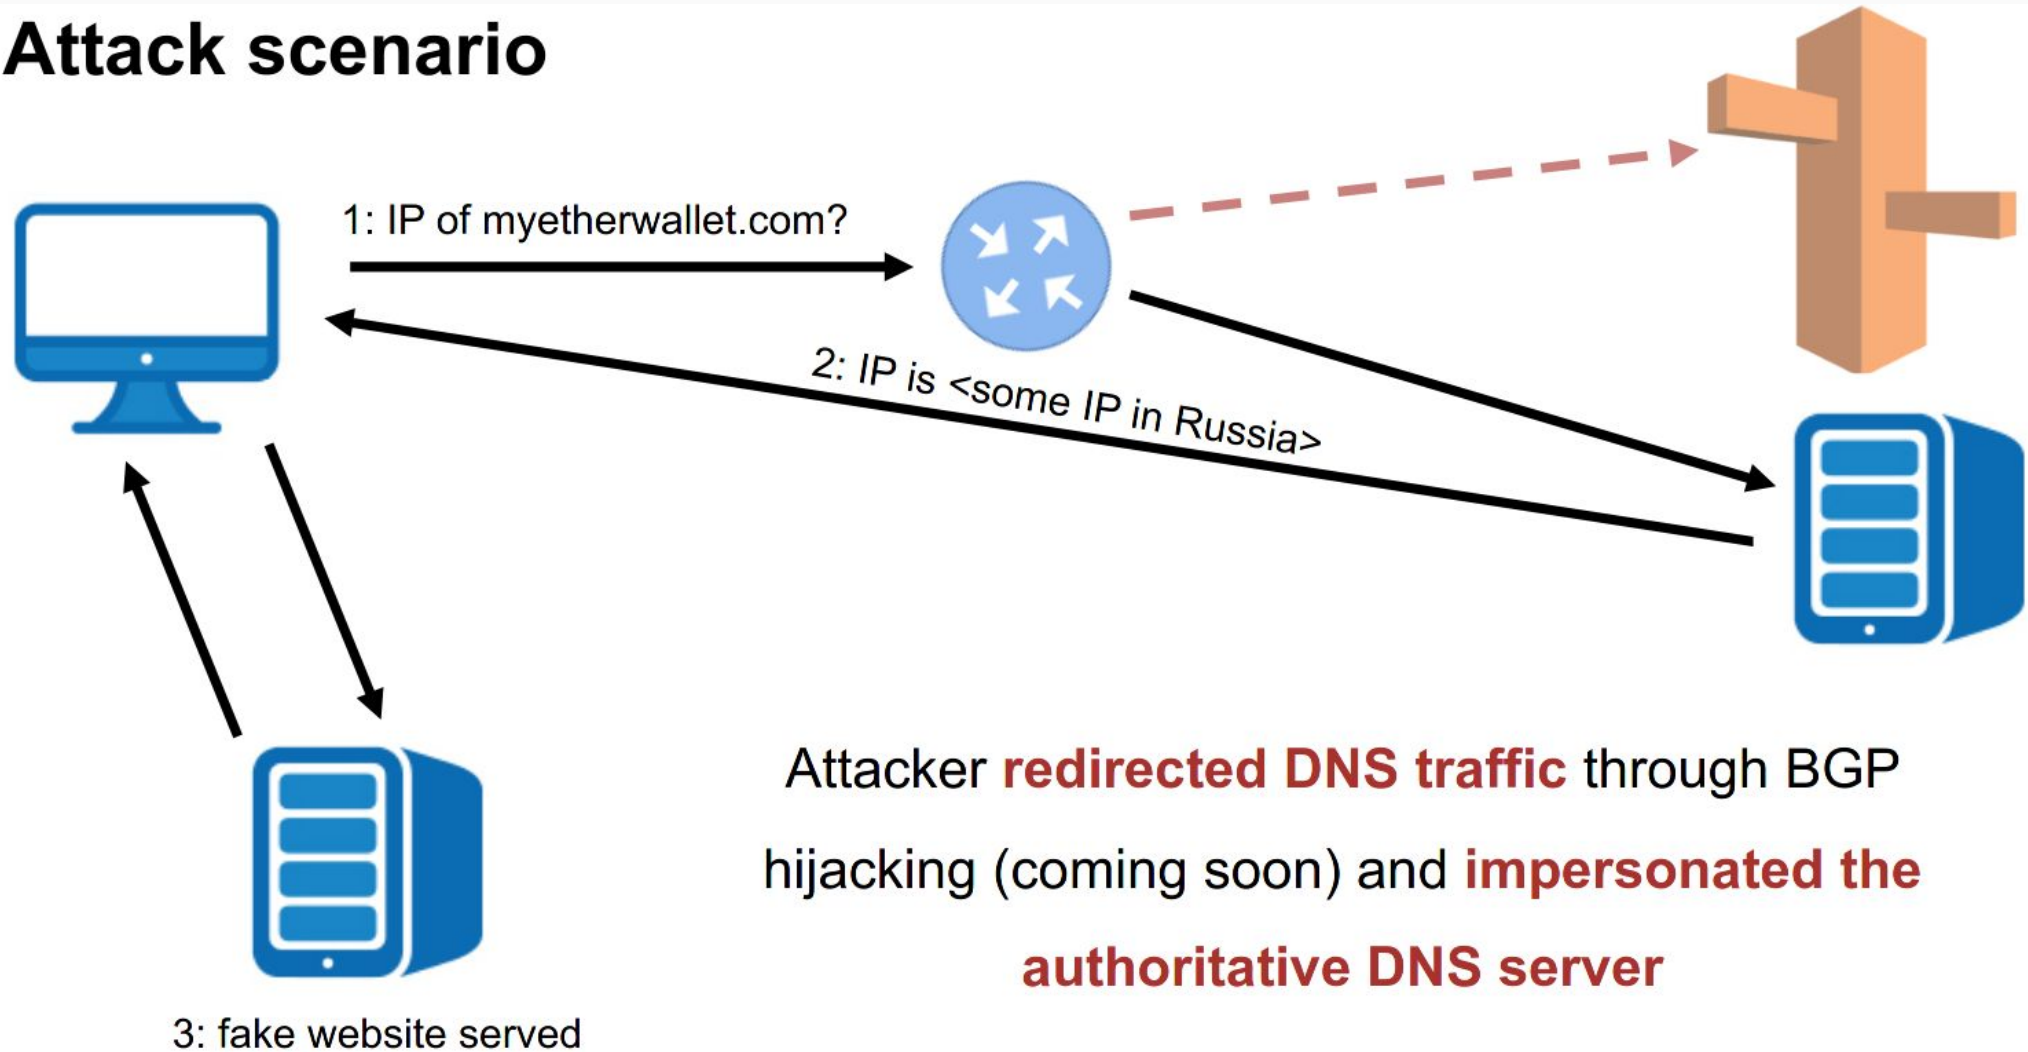
Slide from Tommaso Ciussani

Q3: MyEtherWallet



Q3: MyEtherWallet

Attack scenario



Slide from Tommaso Ciussani

Q3: Hard to fix

1. DNS not authenticated!
2. BGP!
3. Recursive resolvers cached the wrong authoritative answer!



Warning: Potential Security Risk Ahead

Firefox detected a potential security threat and did not continue to wrong.host.badssl.com. If you visit this site, attackers could try to steal information like your passwords, emails, or credit card details.

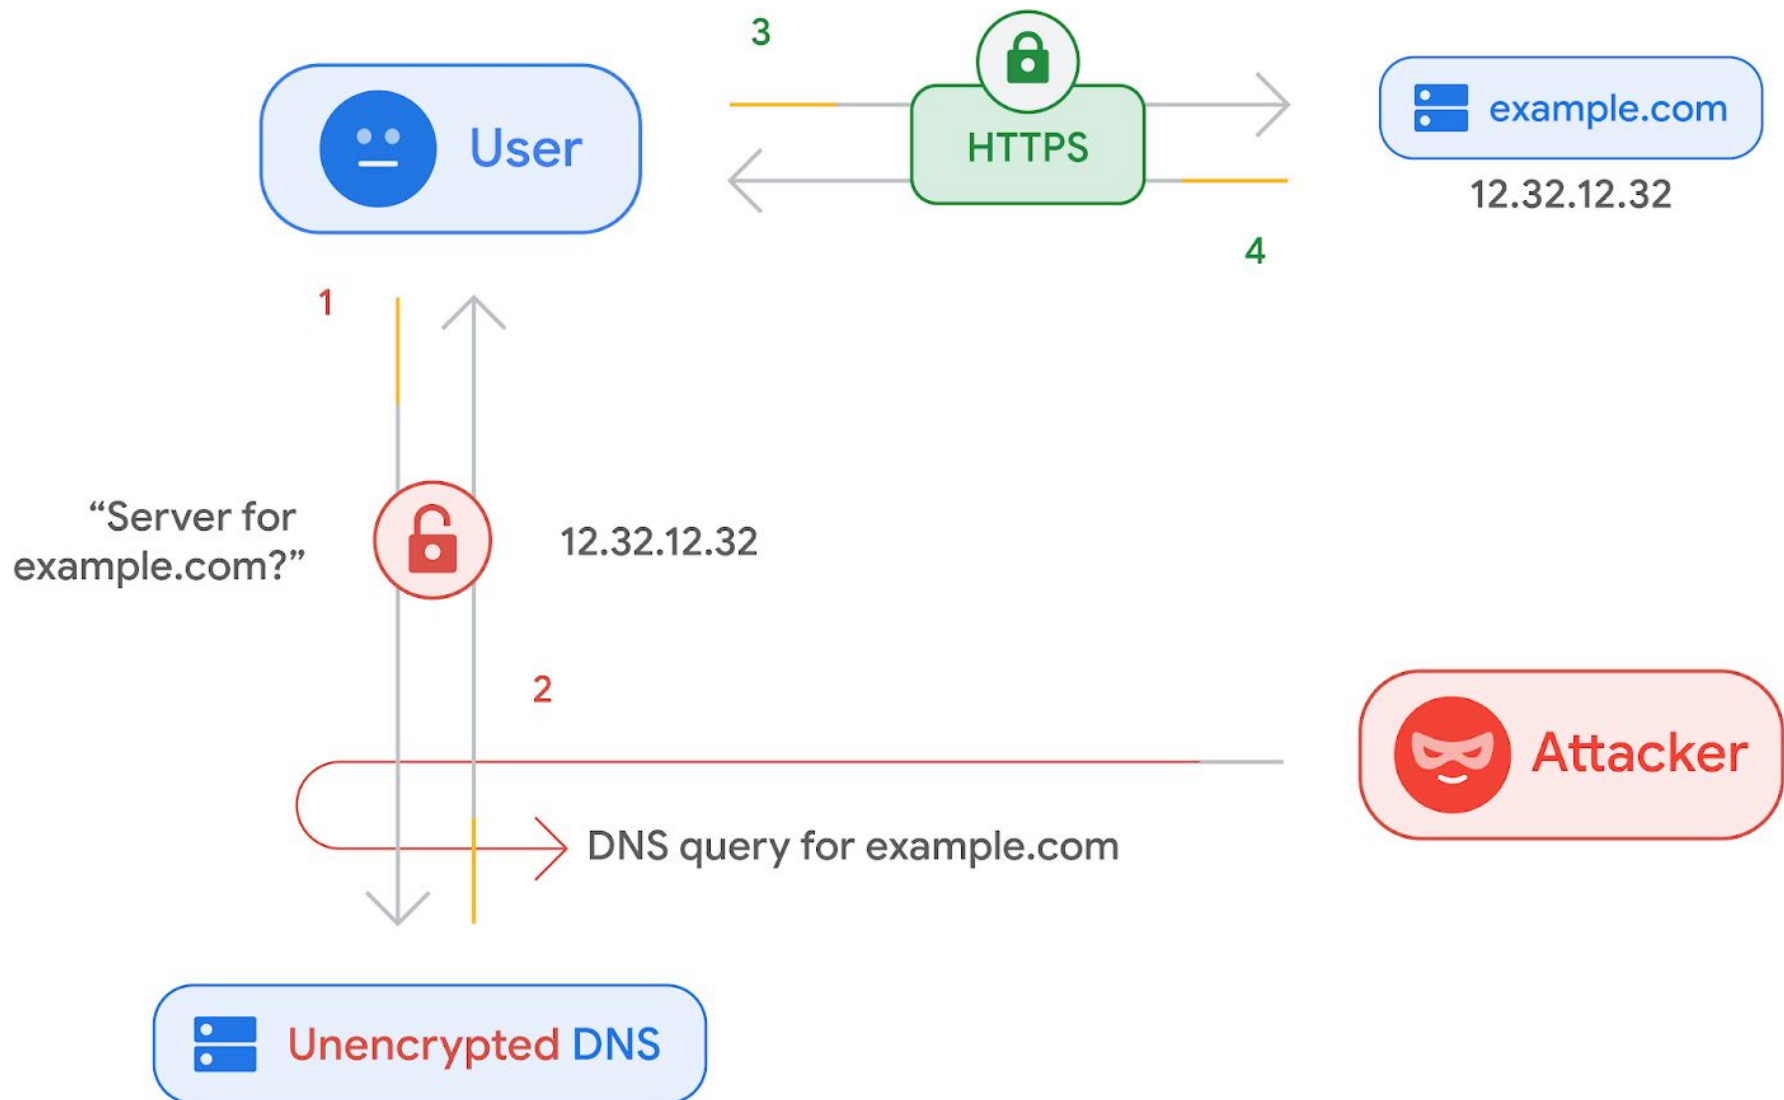
What can you do about it?

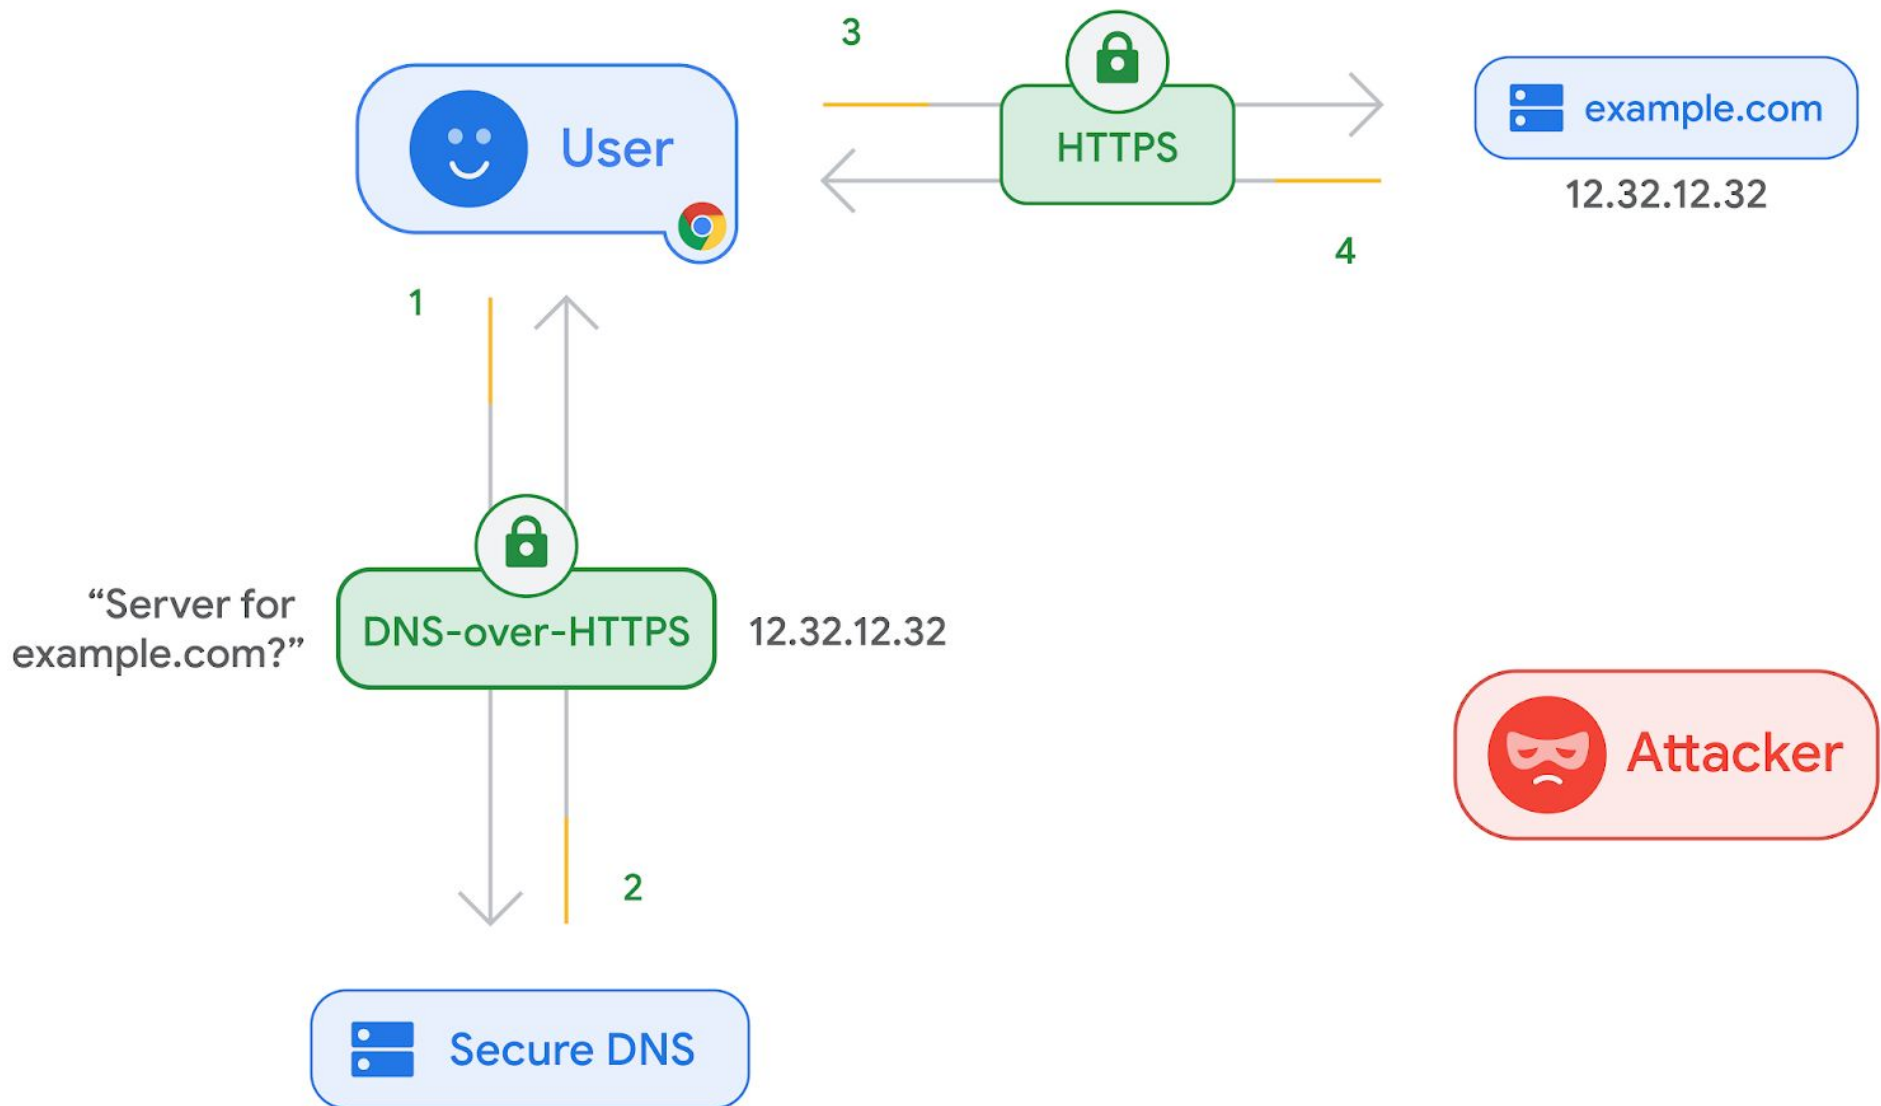
The issue is most likely with the website, and there is nothing you can do to resolve it. You can notify the website's administrator about the problem.

[Learn more...](#)

Go Back (Recommended)

Advanced...









Google Public DNS



```
205     new DohProviderEntry(  
206         "Switch", base::nullopt /* provider_id_for_histogram */,  
207         {"130.59.31.251", "130.59.31.248", "2001:620:0:ff::2",  
208         "2001:620:0:ff::3"},  
209         {"dns.switch.ch"} /* dns_over_tls_hostnames */,  
210         "https://dns.switch.ch/dns-query", "" /* ui_name */,  
211         "" /* privacy_policy */, false /* display_globally */,  
212         {} /* display_countries */),  
213     });  
214     return *providers;  
215 }
```


How does Firefox handle split-horizon DNS?

If Firefox fails to resolve a domain via DoH, it will fall back to the DNS. This means that any domains that are only available on the ordinary DNS (because they aren't public) will be resolved that way. If you have a domain that is publicly resolvable but resolves differently internally, then you should use enterprise settings to disable DoH.

such as family-safe filtering, and therefore avoid breaking user expectations. Furthermore, if there's any hiccup with the DNS-over-HTTPS connection, Chrome will fall back to the regular DNS service of the user's current provider by default, in order to avoid any disruption, while periodically retrying to secure the DNS communication. Finally, to avoid an issue that otherwise could arise